



# Classification Tuning

## XGBoost


### XGBoost (No resampling)



### XGBoost (Undersampling)




### XGBoost (Oversampling)




## Random Forest


### Random Forest (No resampling)



### Random Forest (Undersampling)



### Random Forest (Oversampling)



## Code

```
# LOAD LIBRARIES
# -----
library(tidyverse)
library(readr)
library(xgboost)
library(Matrix)
library(tictoc)
library(furrr)
library(smotefamily)

# IMPORT DATA, RELEVEL FACTOR COLUMNS
# -----
```

```

winequality <- read_csv("data/winequality-all.csv") %>% mutate(type = factor(type,
  levels = c("red", "white")), type01 = as.numeric(ifelse(type ==
  "white", 0, 1)), qualityclass = factor(qualityclass, levels = c("Low",
  "Normal", "High")), under_class = ifelse(qualityclass ==
  "Normal", 0, 1), over_class = ifelse(qualityclass == "Low",
  0, ifelse(qualityclass == "Normal", 1, 2)))
colnames(winequality) <- make.names(names(winequality), unique = TRUE)
summary(winequality)

# RESAMPLING FUNCTIONS
# -----

# FUNCTION FOR UNDERSAMPLING
undersample <- function(train_df, nsample) {
  df_wine_0_ind <- which(train_df$under_class == 0)
  df_wine_1_ind <- which(train_df$under_class == 1)
  pick_0 <- sample(df_wine_0_ind, nsample)
  undersample_wine <- train_df[c(df_wine_1_ind, pick_0), ] #Final Data frame
  undersample_wine <- undersample_wine %>% dplyr::select(qualityclass,
    type, fixed.acidity, volatile.acidity, citric.acid, residual.sugar,
    chlorides, free.sulfur.dioxide, total.sulfur.dioxide,
    density, pH, sulphates, alcohol)
  # table(undersample_wine$under_class) # have just to make
  # sure it's all balancing out how I think
  return(undersample_wine)
}

# FUNCTION FOR OVERSAMPLING
oversample <- function(train_df, k) {

  winequality_low <- filter(winequality, qualityclass %in%
    c("Low"))
  winequality_normal <- filter(winequality, qualityclass %in%
    c("Normal"))
  winequality_high <- filter(winequality, qualityclass %in%
    c("High"))

  wine <- sort(sample(nrow(winequality_normal), nrow(winequality_normal) *
    0.5))
  winequality_norm1 <- winequality_normal[wine, ]
  winequality_norm2 <- winequality_normal[-wine, ]

  wine_LN <- rbind(winequality_low, winequality_norm1)
  wine_HN <- rbind(winequality_high, winequality_norm2)
  SMOTEData1 <- SMOTE(wine_LN[, c("fixed.acidity", "volatile.acidity",
    "citric.acid", "residual.sugar", "chlorides", "free.sulfur.dioxide",
    "total.sulfur.dioxide", "density", "pH", "sulphates",
    "alcohol", "type01")], wine_LN[, "over_class"], K = k,
    dup_size = 0)
  SMOTEData2 <- SMOTE(wine_HN[, c("fixed.acidity", "volatile.acidity",
    "citric.acid", "residual.sugar", "chlorides", "free.sulfur.dioxide",
    "total.sulfur.dioxide", "density", "pH", "sulphates",
    "alcohol", "type01")], wine_HN[, "over_class"], K = k,

```

```

    dup_size = 0)
oversample_df1 <- SMOTEData1$data #Final data frame
oversample_df2 <- SMOTEData2$data #Final data frame
oversample_df <- rbind(oversample_df1, oversample_df2) %>%
  mutate(type01 = round(type01)) %>% mutate(type = as.factor(ifelse(type01 ==
0, "white", "red")), qualityclass = ifelse(class == 0,
"Low", ifelse(class == "1", "Normal", "High"))) %>% mutate(qualityclass = factor(qualityclass,
levels = c("Low", "Normal", "High"))) %>% dplyr::select(qualityclass,
type, fixed.acidity, volatile.acidity, citric.acid, residual.sugar,
chlorides, free.sulfur.dioxide, total.sulfur.dioxide,
density, pH, sulphates, alcohol)
# table(oversample_df$class) have just to make sure it's all
# balancing out how I think
return(oversample_df)
}

# XGBOOST FUNCTIONS
# -----

# SET UP FUNCTION TO EVALUATE XGBOOST
xgbFunc <- function(df = winequality, samplingMethod = "none",
nUndersample = 2000, kOversample = 5, trainPct = 0.7, max.depth,
eta, nround = 2, nthread = 2, show.table = F) {

  require(Matrix)
  require(xgboost)

  # set up training/testing sets
  n <- nrow(df)
  train.index <- sample(seq(1, n), floor(n * trainPct), replace = F)

  # create dgCMatrx for modeling

  # training
  train.data <- df[train.index, ] # Normal

  if (samplingMethod == "undersample") {
    train.data <- undersample(train.data, nUndersample) # Undersample
  }

  if (samplingMethod == "oversample") {
    train.data <- oversample(train.data, kOversample) # Oversample
  }

  train.data <- train.data %>% dplyr::select(qualityclass,
type, fixed.acidity, volatile.acidity, citric.acid, residual.sugar,
chlorides, free.sulfur.dioxide, total.sulfur.dioxide,
density, pH, sulphates, alcohol)

  train.datamatrix <- sparse.model.matrix(qualityclass ~ .,
data = train.data)[, -1]
  train.qualityclass <- train.data$qualityclass
  train.label <- as.integer(train.data$qualityclass) - 1 # label conversion

```

```

xgb.train <- list(data = train.datamatrix, label = train.label)

# testing
test.data <- df[-train.index, ] %>% dplyr::select(qualityclass,
  type, fixed.acidity, volatile.acidity, citric.acid, residual.sugar,
  chlorides, free.sulfur.dioxide, total.sulfur.dioxide,
  density, pH, sulphates, alcohol)
test.datamatrix <- sparse.model.matrix(qualityclass ~ .,
  data = test.data)[, -1]
test.qualityclass <- test.data$qualityclass
test.label <- as.integer(test.data$qualityclass) - 1 # label conversion
xgb.test <- list(data = test.datamatrix, label = test.label)

# fit xgboost model
xgb.fit <- xgboost(data = xgb.train$data, label = xgb.train$label,
  booster = "gbtree", max.depth = max.depth, eta = eta,
  nround = nround, nthread = nthread, objective = "multi:softprob",
  eval_metric = "merror", num_class = length(levels(train.qualityclass)),
  verbose = 0)

# predict
xgb.pred = predict(xgb.fit, xgb.test$data, reshape = T) %>%
  as.data.frame()
colnames(xgb.pred) = levels(train.qualityclass)
xgb.pred$prediction = apply(xgb.pred, 1, function(x) colnames(xgb.pred)[which.max(x)])
xgb.pred$label = levels(train.qualityclass)[test.label +
  1]
xgb.pred <- xgb.pred %>% mutate(prediction = factor(prediction,
  levels = c("Low", "Normal", "High")), label = factor(label,
  levels = c("Low", "Normal", "High")))

# evaluated prediction
accuracy.all <- mean(xgb.pred$prediction == xgb.pred$label)
table <- with(xgb.pred, table(label, prediction))
prop.table <- table/rowSums(table)
accuracy.low <- prop.table[1, 1]
accuracy.normal <- prop.table[2, 2]
accuracy.high <- prop.table[3, 3]
accuracy <- cbind(accuracy.all, accuracy.low, accuracy.normal,
  accuracy.high)

if (show.table) {
  return(list(accuracy = accuracy, table = table, prop.table = prop.table))
} else {
  return(accuracy)
}
}

# XGBOOST MCMC FUNCTION WITH PARALLEL COMPUTING
xgbMCMC <- function(samplingMethod = "none", nUndersample = 2000,
  kOversample = 5, B = 5, trainPct = 0.7, max.depth, eta, nround = 2,

```

```

nthread = 2) {
  require(furrr)

  # Create Parameter Grid
  mcmc.grid <- expand_grid(B = seq(1, B), samplingMethod = samplingMethod,
    nUndersample = nUndersample, kOversample = kOversample,
    trainPct = trainPct, max.depth = max.depth, eta = eta,
    nround = nround, nthread = nthread)

  # Obtain Accuracy
  accuracyList <- furrr::future_pmap(mcmc.grid[, -1], xgbFunc)
  xgbAccuracy <- matrix(unlist(accuracyList, use.names = TRUE),
    ncol = 4, nrow = nrow(mcmc.grid), byrow = T)
  colnames(xgbAccuracy) <- colnames(accuracyList[[1]])

  # Summarize Accuracy
  results <- cbind(mcmc.grid, xgbAccuracy) %>% pivot_longer(cols = c("accuracy.all",
    "accuracy.low", "accuracy.normal", "accuracy.high"),
    names_to = "accuracyGroup", values_to = "accuracy") %>%
    mutate(Method = "XGBoost") %>% dplyr::group_by(Method,
    accuracyGroup, samplingMethod, nUndersample, kOversample,
    max.depth, eta, nround, nthread) %>% summarise(B = n(),
    mean = mean(accuracy), lower = quantile(accuracy, probs = c(0.05)),
    upper = quantile(accuracy, probs = c(0.95))) %>% ungroup()

  return(results)
}

# RANDOM FOREST FUNCTIONS
# -----

# SET UP FUNCTION TO EVALUATE RANDOM FOREST
rfFunc <- function(df = winequality, samplingMethod = "none",
  nUndersample = 2000, kOversample = 5, trainPct = 0.7, importance = F,
  mtry = 4, ntree = 500, show.table = F) {

  require(randomForest)

  # set up training/testing sets
  n <- nrow(df)
  train.index <- sample(seq(1, n), floor(n * trainPct), replace = F)

  # training
  train.data <- df[train.index, ] # Normal

  if (samplingMethod == "undersample") {
    train.data <- undersample(train.data, nUndersample)
  }

  if (samplingMethod == "oversample") {
    train.data <- oversample(train.data, kOversample)
  }
}

```

```

train.data <- train.data %>% dplyr::select(qualityclass,
  type, fixed.acidity, volatile.acidity, citric.acid, residual.sugar,
  chlorides, free.sulfur.dioxide, total.sulfur.dioxide,
  density, pH, sulphates, alcohol)

# testing
test.data <- df[-train.index, ] %>% dplyr::select(qualityclass,
  type, fixed.acidity, volatile.acidity, citric.acid, residual.sugar,
  chlorides, free.sulfur.dioxide, total.sulfur.dioxide,
  density, pH, sulphates, alcohol)

# Fit that Random Forest!
rf.fit <- randomForest(qualityclass ~ ., data = train.data,
  method = "class", ntree = ntree, mtry = mtry, importance = importance)

# Get that Prediction!
rf.pred = predict(rf.fit, newdata = test.data)

# Evaluate Prediction
accuracy.all <- mean(rf.pred == test.data$qualityclass) #this is not working SJA
table <- table(test.data$qualityclass, rf.pred)
prop.table <- table/rowSums(table)
accuracy.low <- prop.table[1, 1]
accuracy.normal <- prop.table[2, 2]
accuracy.high <- prop.table[3, 3]
accuracy <- cbind(accuracy.all, accuracy.low, accuracy.normal,
  accuracy.high)

if (show.table) {
  return(list(accuracy = accuracy, table = table, prop.table = prop.table))
} else {
  return(accuracy)
}
}

# RANDOM FOREST MCMC FUNCTION WITH PARALLEL COMPUTING
rfMCMC <- function(B = 5, samplingMethod = "none", nUndersample = 2000,
  kOversample = 5, trainPct = 0.7, ntree = 500, mtry = 4) {
  require(furrr)

  # Create Parameter Grid
  mcmc.grid <- expand_grid(B = seq(1, B), samplingMethod = samplingMethod,
    nUndersample = nUndersample, kOversample = kOversample,
    trainPct = trainPct, ntree = ntree, mtry = mtry)

  # Obtain Accuracy
  accuracyList <- furrr::future_pmap(mcmc.grid[, -1], rfFunc)
  rfAccuracy <- matrix(unlist(accuracyList, use.names = TRUE),
    ncol = 4, nrow = nrow(mcmc.grid), byrow = T)
  colnames(rfAccuracy) <- colnames(accuracyList[[1]])

  # Summarize Accuracy

```

```

results <- cbind(mcmc.grid, rfAccuracy) %>% pivot_longer(cols = c("accuracy.all",
  "accuracy.low", "accuracy.normal", "accuracy.high"),
  names_to = "accuracyGroup", values_to = "accuracy") %>%
mutate(Method = "Random Forest") %>% dplyr::group_by(Method,
  accuracyGroup, samplingMethod, nUndersample, kOversample,
  ntree, mtry) %>% summarise(B = n(), mean = mean(accuracy),
  lower = quantile(accuracy, probs = c(0.05)), upper = quantile(accuracy,
  probs = c(0.95))) %>% ungroup()

return(results)
}

# HYPERPARAMETER GRID SEARCH
# -----
setB = 50

# XGBOOST

tic()
xgbMCMC.none.gridsearch <- xgbMCMC(samplingMethod = "none", nUndersample = NA,
  kOversample = NA, trainPct = 0.7, B = setB, max.depth = seq(50,
  250, 50), eta = c(0.05, 0.1, 0.25, 0.5, 0.75, 0.9, 0.95),
  nround = seq(2, 12, 2), nthread = seq(2, 12, 2))
toc()

xgbMCMC.none.gridsearch <- xgbMCMC.none.gridsearch %>% filter(accuracyGroup ==
  "accuracy.all") %>% arrange(-mean)
xgbMCMC.none.gridsearch
# write.csv(xgbMCMC.none.gridsearch, file =
# 'reports/xgbMCMC.none.gridsearch.csv', row.names = F, na =
# '')

xgbMCMC.none.gridsearch.plot <- xgbMCMC.none.gridsearch[c(0:10),
  ] %>% mutate(label = paste("Max Depth = ", max.depth, "; Eta = ",
  eta, "; nround = ", nround, "; nthread = ", nthread, sep = "")) %>%
ggplot(aes(x = mean, y = reorder(label, mean))) + geom_point() +
geom_errorbar(aes(xmin = lower, xmax = upper)) + theme_bw() +
scale_y_discrete("") + scale_x_continuous("Overall Accuracy",
  labels = scales::percent) + ggtitle("XGBoost (No resampling)")
xgbMCMC.none.gridsearch.plot

tic()
xgbMCMC.undersample.gridsearch <- xgbMCMC(samplingMethod = "undersample",
  nUndersample = seq(1000, 2000, 500), kOversample = NA, trainPct = 0.7,
  B = setB, max.depth = seq(50, 250, 50), eta = c(0.5, 0.1,
  0.25, 0.5, 0.75, 0.9, 0.95), nround = seq(2, 12, 2),
  nthread = seq(2, 12, 2))
toc()
xgbMCMC.undersample.gridsearch <- xgbMCMC.undersample.gridsearch %>%
  select(accuracyGroup == "accuracy.all") %>% arrange(-mean)
xgbMCMC.undersample.gridsearch

```



```

# write.csv(xgbMCMC.undersample.gridsearch, file =
# 'reports/xgbMCMC.undersample.gridsearch.csv', row.names =
# F, na = '')

xgbMCMC.undersample.gridsearch.plot <- xgbMCMC.undersample.gridsearch[c(0:10),
] %>% mutate(label = paste("Max Depth = ", max.depth, "; Eta = ",
eta, "; nround = ", nround, "; nthread = ", nthread, " (N = ",
nundersample, ")", sep = ""))
ggplot(aes(x = mean, y = reorder(label, mean))) + geom_point() +
geom_errorbar(aes(xmin = lower, xmax = upper)) + theme_bw() +
scale_y_discrete("") + scale_x_continuous("Overall Accuracy",
labels = scales::percent) + ggtitle("XGBoost (Undersampling)")
xgbMCMC.undersample.gridsearch.plot

tic()
xgbMCMC.oversample.gridsearch <- xgbMCMC(samplingMethod = "oversample",
nUndersample = NA, kOversample = seq(3, 7, 2), trainPct = 0.7,
B = setB, max.depth = seq(50, 250, 50), eta = c(0.5, 0.1,
0.25, 0.5, 0.75, 0.9, 0.95), nround = seq(2, 12, 2),
nthread = seq(2, 12, 2))
toc()
xgbMCMC.oversample.gridsearch <- xgbMCMC.oversample.gridsearch %>%
select(accuracyGroup == "accuracy.all") %>% arrange(-mean)
xgbMCMC.oversample.gridsearch
# write.csv(xgbMCMC.oversample.gridsearch, file =
# 'reports/xgbMCMC.oversample.gridsearch.csv', row.names = F,
# na = '')

xgbMCMC.oversample.gridsearch.plot <- xgbMCMC.oversample.gridsearch[c(0:10),
] %>% mutate(label = paste("Max Depth = ", max.depth, "; Eta = ",
eta, "; nround = ", nround, "; nthread = ", nthread, " (k = ",
kOversample, ")", sep = ""))
ggplot(aes(x = mean, y = reorder(label, mean))) + geom_point() +
geom_errorbar(aes(xmin = lower, xmax = upper)) + theme_bw() +
scale_y_discrete("") + scale_x_continuous("Overall Accuracy",
labels = scales::percent) + ggtitle("XGBoost (Oversampling)")
xgbMCMC.oversample.gridsearch.plot

# RANDOM FOREST

# No resampling
tic()
rfMCMC.none.gridsearch <- rfMCMC(samplingMethod = "none", nUndersample = NA,
kOversample = NA, trainPct = 0.7, B = setB, ntree = seq(200,
800, 200), mtry = seq(2, 10, 2) # Default should be 4
)
toc()

rfMCMC.none.gridsearch <- rfMCMC.none.gridsearch %>% filter(accuracyGroup ==
"accuracy.all") %>% arrange(-mean)
rfMCMC.none.gridsearch
# write.csv(rfMCMC.none.gridsearch, file =
# 'reports/rfMCMC.none.gridsearch.csv', row.names = F, na =

```

```

# '')

rfMCMC.none.gridsearch.plot <- rfMCMC.none.gridsearch[c(0:10),
] %>% mutate(label = paste("Ntree = ", ntree, "; mtry = ",
mtry, sep = "")) %>% ggplot(aes(x = mean, y = reorder(label,
mean))) + geom_point() + geom_errorbar(aes(xmin = lower,
xmax = upper)) + theme_bw() + scale_y_discrete("") + scale_x_continuous("Overall Accuracy",
labels = scales::percent) + ggtitle("Random Forest (No resampling)")
rfMCMC.none.gridsearch.plot

# Undersample
tic()
rfMCMC.undersample.gridsearch <- rfMCMC(samplingMethod = "undersample",
nUndersample = seq(1000, 2000, 500), kOversample = NA, trainPct = 0.7,
B = setB, ntree = seq(200, 800, 200), mtry = seq(2, 10, 2) # Default should be 4
)
toc()

rfMCMC.undersample.gridsearch <- rfMCMC.undersample.gridsearch %>%
filter(accuracyGroup == "accuracy.all") %>% arrange(-mean)
rfMCMC.undersample.gridsearch
# write.csv(rfMCMC.undersample.gridsearch, file =
# 'reports/rfMCMC.undersample.gridsearch.csv', row.names = F,
# na = '')

rfMCMC.undersample.gridsearch.plot <- rfMCMC.undersample.gridsearch[c(0:10),
] %>% mutate(label = paste("Ntree = ", ntree, "; mtry = ",
mtry, " (N = ", nUndersample, ")", sep = "")) %>% ggplot(aes(x = mean,
y = reorder(label, mean))) + geom_point() + geom_errorbar(aes(xmin = lower,
xmax = upper)) + theme_bw() + scale_y_discrete("") + scale_x_continuous("Overall Accuracy",
labels = scales::percent) + ggtitle("Random Forest (Undersampling)")
rfMCMC.undersample.gridsearch.plot

# Oversample
tic()
rfMCMC.oversample.gridsearch <- rfMCMC(samplingMethod = "oversample",
nUndersample = NA, kOversample = seq(3, 7, 2), trainPct = 0.7,
B = setB, ntree = seq(200, 800, 200), mtry = seq(2, 10, 2) # Default should be 4
)
toc()

rfMCMC.oversample.gridsearch <- rfMCMC.oversample.gridsearch %>%
filter(accuracyGroup == "accuracy.all") %>% arrange(-mean)
rfMCMC.oversample.gridsearch
# write.csv(rfMCMC.oversample.gridsearch, file =
# 'reports/rfMCMC.oversample.gridsearch.csv', row.names = F,
# na = '')

rfMCMC.oversample.gridsearch.plot <- rfMCMC.oversample.gridsearch[c(0:10),
] %>% mutate(label = paste("Ntree = ", ntree, "; mtry = ",
mtry, " (k = ", kOversample, ")", sep = "")) %>% ggplot(aes(x = mean,
y = reorder(label, mean))) + geom_point() + geom_errorbar(aes(xmin = lower,
xmax = upper)) + theme_bw() + scale_y_discrete("") + scale_x_continuous("Overall Accuracy",

```

```

    labels = scales::percent) + ggtitle("Random Forest (Oversampling)")
rfMCMC.oversample.gridsearch.plot

# FINAL RESULTS
# -----

# XGBOOST MCMC RESULTS
setB = 100

tic()
xgbMCMC.none.results <- xgbMCMC(samplingMethod = "none", nUndersample = NA,
    kOversample = NA, trainPct = 0.7, B = setB, max.depth = 50,
    eta = 0.75, nround = 12, nthread = 8)
toc()

tic()
xgbMCMC.undersample.results <- xgbMCMC(samplingMethod = "undersample",
    nUndersample = 1000, kOversample = NA, trainPct = 0.7, B = setB,
    max.depth = 50, eta = 0.1, nround = 2, nthread = 2)
toc()

tic()
xgbMCMC.oversample.results <- xgbMCMC(samplingMethod = "oversample",
    nUndersample = NA, kOversample = 3, trainPct = 0.7, B = setB,
    max.depth = 50, eta = 0.1, nround = 2, nthread = 2)
toc()

xgbMCMC.results <- rbind(xgbMCMC.none.results, xgbMCMC.undersample.results,
    xgbMCMC.oversample.results) %>% mutate(label = paste("Max Depth = ",
    max.depth, "; Eta = ", eta, sep = "")) %>% select(-max.depth,
    -eta, -nround, -nthread)

# RANDOM FORST MCMC RESULTS
tic()
rfMCMC.none.results <- rfMCMC(B = setB, samplingMethod = "none",
    nUndersample = NA, kOversample = NA, trainPct = 0.7, ntree = 200,
    mtry = 2)
toc()

tic()
rfMCMC.undersample.results <- rfMCMC(B = setB, samplingMethod = "undersample",
    nUndersample = 2000, kOversample = NA, trainPct = 0.7, ntree = 800,
    mtry = 2)
toc()

tic()
rfMCMC.oversample.results <- rfMCMC(B = setB, samplingMethod = "oversample",
    nUndersample = NA, kOversample = 3, trainPct = 0.7, ntree = 200,
    mtry = 3)
toc()

```

```

rfMCMC.results <- rbind(rfMCMC.none.results, rfMCMC.undersample.results,
  rfMCMC.oversample.results) %>% mutate(label = paste("N Trees = ",
  ntree, "; Mtry = ", mtry, sep = "")) %>% select(-ntree, -mtry)

# COMBINE XGBOOST AND RANDOM FOREST RESULTS & PLOT
MCMC.results <- rbind(xgbMCMC.results, rfMCMC.results)

write.csv(MCMC.results, "reports/MCMC.results.csv", row.names = F,
  na = "")
MCMC.results <- read_csv("reports/MCMC.results.csv")

# new facet labels
accuracyGroup.labs <- c("Overall", "Low Quality", "Normal Quality",
  "High Quality")
names(accuracyGroup.labs) <- c("accuracy.all", "accuracy.low",
  "accuracy.normal", "accuracy.high")

# plot accuracy
MCMC.results %>% mutate(samplingMethod = factor(samplingMethod,
  levels = c("none", "undersample", "oversample")), accuracyGroup = factor(accuracyGroup,
  levels = c("accuracy.all", "accuracy.low", "accuracy.normal",
  "accuracy.high"))) %>% ggplot(aes(x = mean, y = samplingMethod,
  group = Method, color = Method)) + geom_point(size = 1, position = position_dodge(width = 0.5)) +
  geom_errorbar(aes(xmin = lower, xmax = upper), position = position_dodge(width = 0.5),
  width = 0.4) + facet_wrap(~accuracyGroup, ncol = 1, labeller = labeller(accuracyGroup = accuracyGroup.labs)) +
  theme_bw() + theme(aspect.ratio = 0.2) + scale_y_discrete("Resampling Method") +
  scale_x_continuous("Accuracy", limits = c(0, 1), breaks = seq(0,
  1, 0.2), labels = scales::percent) + scale_color_brewer("",
  palette = "Paired")

```